


# How much annotation time do you get back?

New time  $\approx$  lookup + (1 – cache hit rate)  $\times$  your current annotation time


## Find your expected cache reuse

Dark = annotation work left · green = waiting time returned




## Fast or slow pipeline?

At 80% reuse, the same fraction of expensive annotation work disappears



## Ten samples or one thousand: savings add up

Example: 1-hour pipeline, 80% reuse, bundled cache with no build cost



## ✓ Same annotations

141/141 benchmark samples pass output-equivalence checks

Any unexpected annotation change fails the task.