


Alternative view: testing the runtime-scaling model

Panels, exomes and genomes place the same cache-hit model on very different absolute time scales

A

The miss fraction predicts work remaining

Dashed line is the zero-overhead runtime model




- 1000G-derived panel
- 1000G WGS
- SGDP WGS
- External WGS
- Source-overlap calibration
- 1000G-derived WES
- KPGP WGS
- PGP WGS

B

Lookup overhead matters most for short pipelines

Inferred overhead = measured cached time – miss fraction × direct time



C

How much of the ideal speedup is realized?

100% means zero lookup and preprocessing overhead

