VCFcache converts annotation complexity into time saved

One 4.80-million-variant real WGS · bundled Zenodo gnomAD AF ≥ 10% cache · ~80% hits · light statistics · no technical repeats

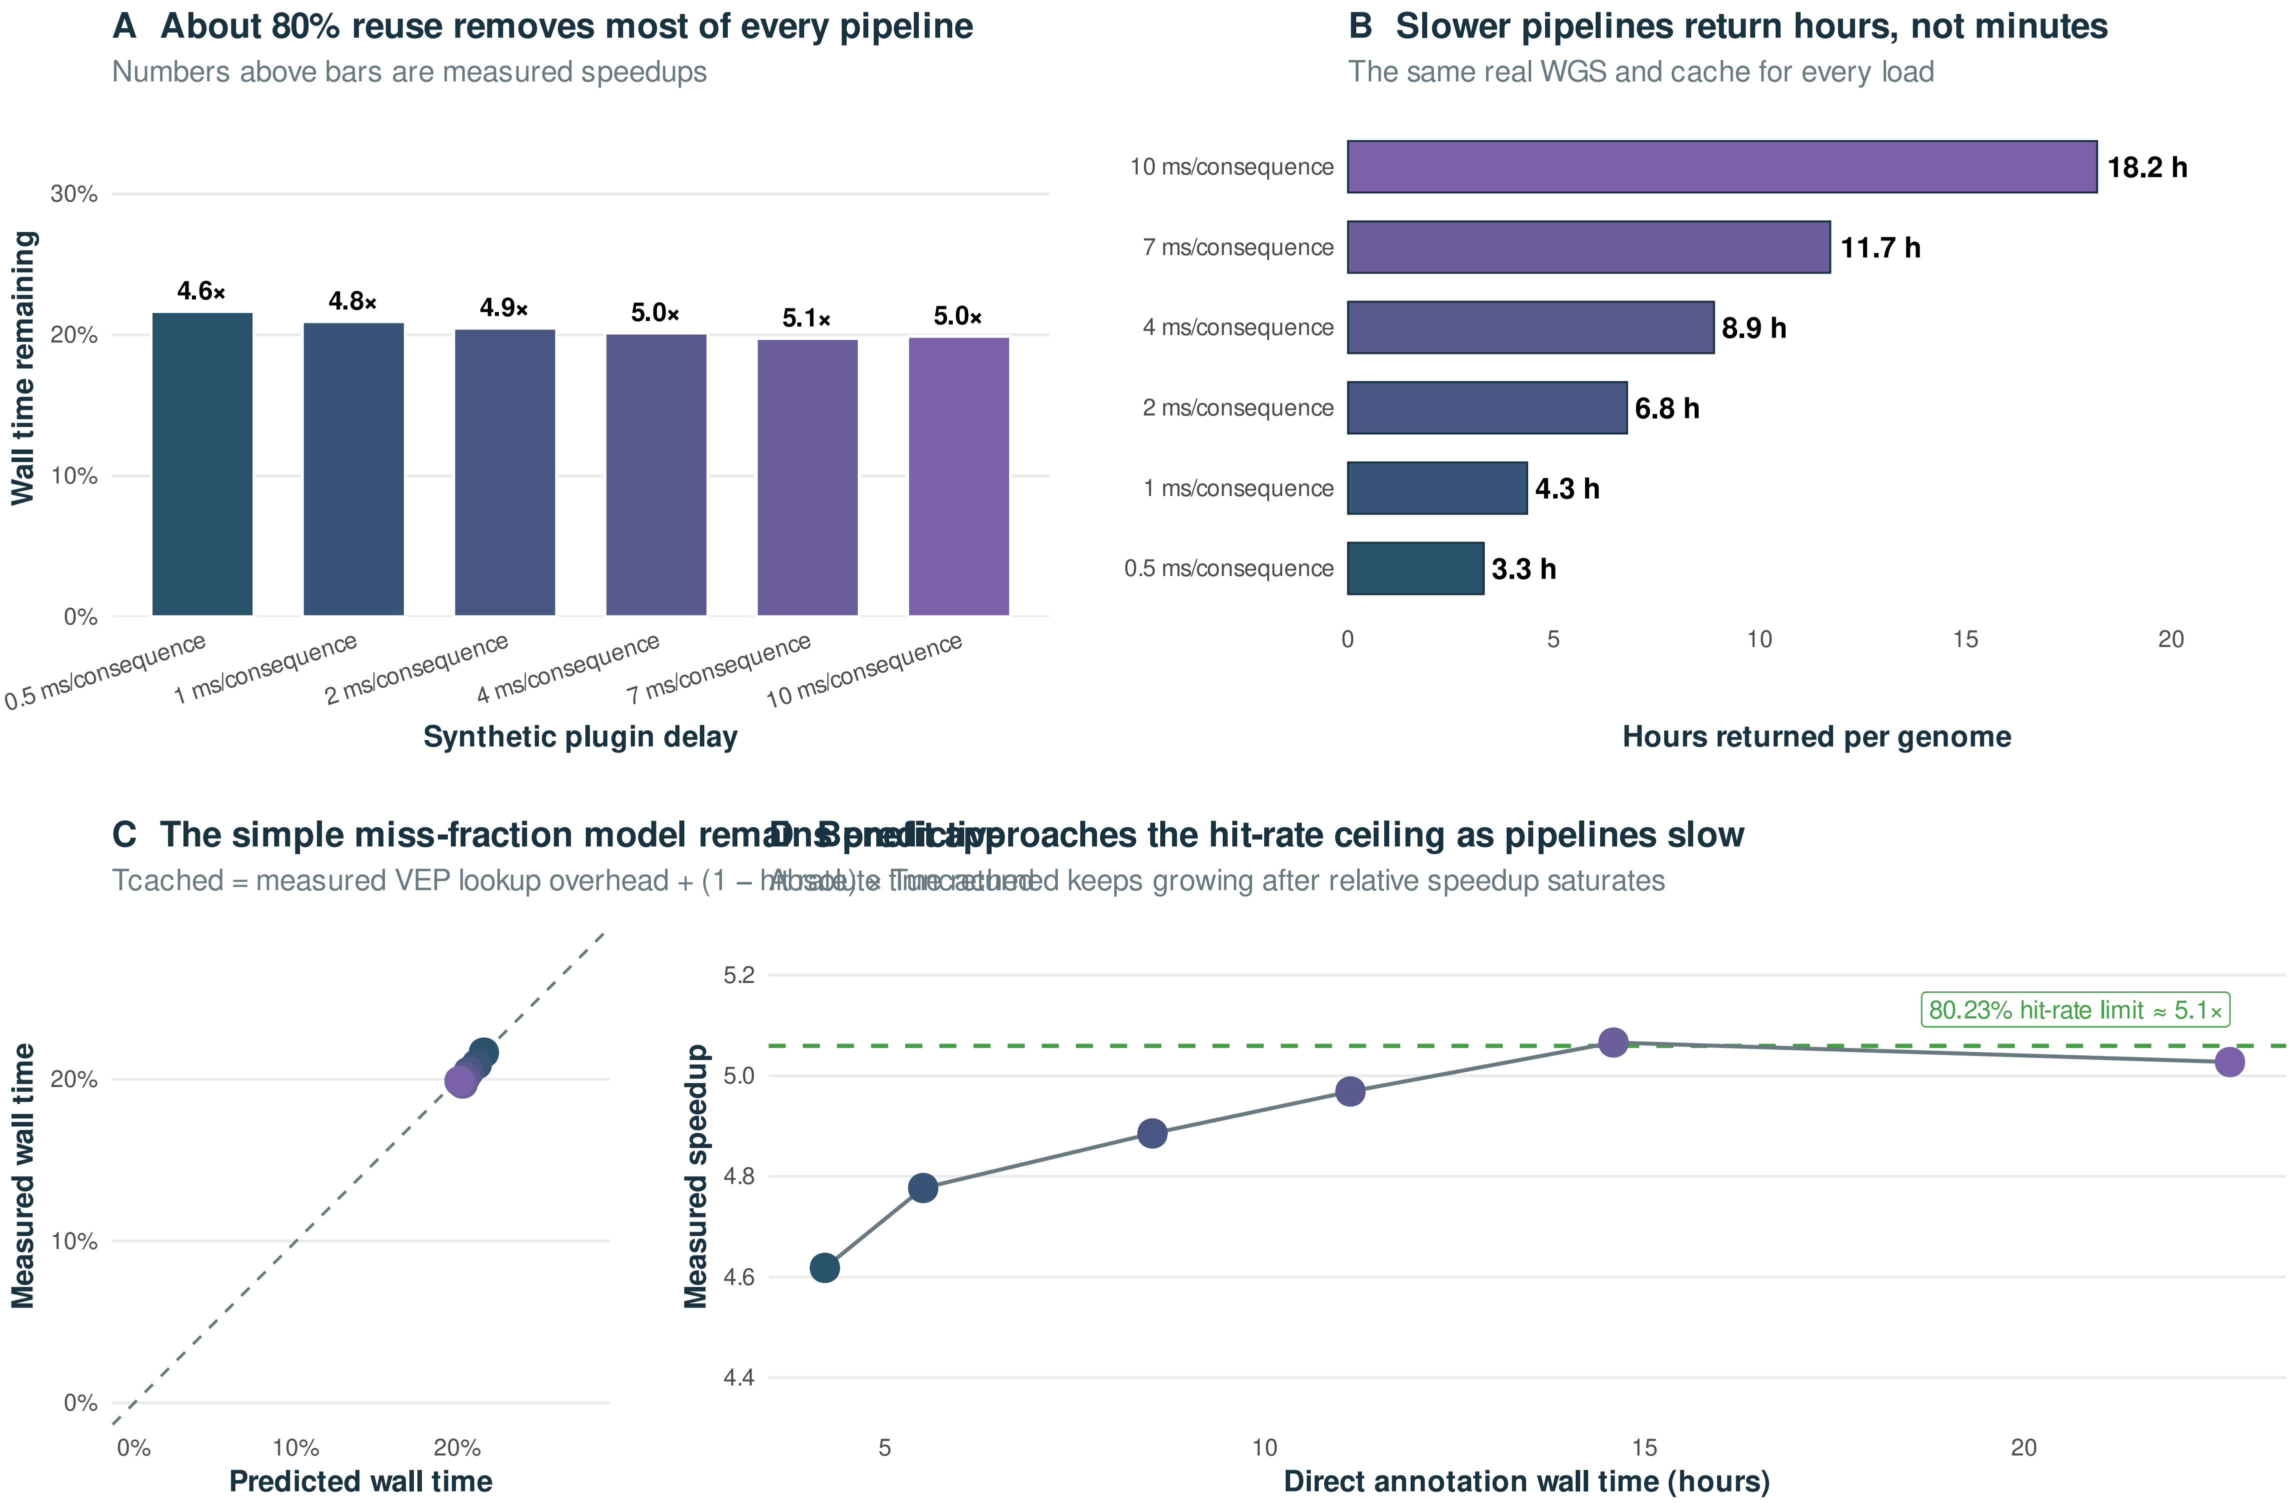

All six cached outputs were semantically equivalent to their pipeline-specific uncached output. Synthetic delays run only for transcript consequences submitted to VEP and emit no annotation fields.