


Alternative view: testing the runtime-scaling model

Panels, exomes and genomes place the same cache-hit model on very different absolute time scales

A

The miss fraction predicts work remaining

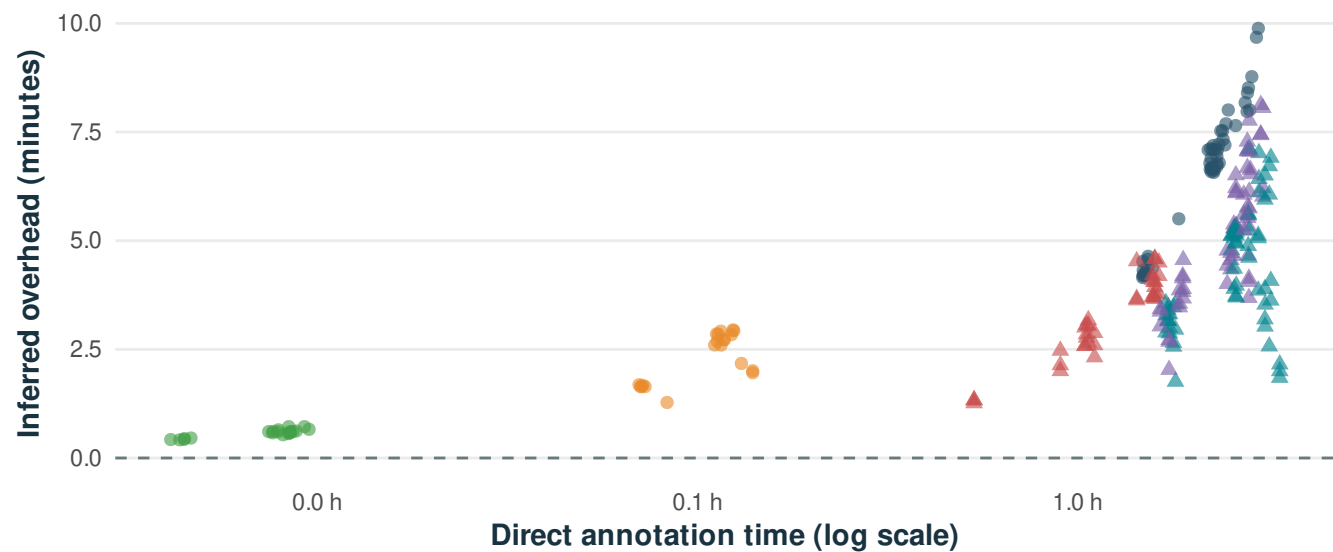
Dashed line is the zero-overhead runtime model



B

Lookup overhead matters most for short pipelines

Inferred overhead = measured cached time – miss fraction × direct time



C

How much of the ideal speedup is realized?

100% means zero lookup and preprocessing overhead

